

LATENCY IN MULTI-REGION DEPLOYMENTS

LATENCY IN DATABASES

How fast can I
read
my data?



How fast can I
write
my data?



SOME REASONS FOR DATABASE LATENCY


- Slow executing SQL statements.
- Speed of network latency, directly related to **where** your data is located
 - Time it takes for read-data to travel to users.
 - Time it takes to write data *and reach consensus* .



COCKROACHDB LOCALITY SETTINGS ADDRESS SPEED-OF-NETWORK LATENCY.




MULTI-REGION LOCALITY OVERVIEW



- In CockroachDB, latency is tuned using multi-region `locality` settings.
- Locality is configured at three levels - cluster, database, table.



TABLE LOCALITY




- **Regional tables** - optimized for fast reads and writes from the *home region*.
- **Regional by row tables** - optimized for fast reads and writes from the *region specified at the row level*.
- **Global tables** - optimized for fast reads from *all regions*.




COCKROACHDB LOCALITY TERMINOLOGY

Region

Cluster



Availability Zone




Locality




REGION -- LOGICAL IDENTIFICATION OF HOW
NODES AND DATA ARE CLUSTERED AROUND
GEOGRAPHICAL LOCATIONS



COCKROACHDB CLUSTER REGION




COCKROACHDB AVAILABILITY ZONE



- A zone is a deployment area within a cluster region.
 - A cluster region is made up of a collection of zones.
- There can be multiple nodes in a single availability zone
 - *Recommendation:* place different replicas of your data in different availability zones.



COCKROACHDB DATABASE REGION



**YOUR MULTI-REGION JOURNEY STARTS WITH
CONFIGURING YOUR CLUSTER REGIONS**



COCKROACHDB CLUSTER REGIONS

- Run CockroachDB on specified infrastructure (physical nodes)
- Define the locality of the nodes available for the databases to use



COCKROACHDB CLUSTER REGIONS

- Specifying a region with a region tier is required in order to enable CockroachDB's multi-region capabilities.
- CockroachDB spreads the replicas of each piece of data across as diverse a set of localities as possible, with the order determining the priority.



COCKROACHDB CLUSTER REGION LIMITATIONS

- Every node requires the same `locality` key-pair settings.
- In self-hosted mode, region locality is not be guaranteed.
 - CockroachDB does not check that a node's physical location matches its user-given locality label.
- You cannot add regions to a CockroachDB Dedicated cluster on the fly.



CONFIGURE CLUSTER REGIONS



- Steps to configuring a cluster region depend on your environment



CONFIGURING CLUSTER REGION

- With demo use `cockroach demo`
- With CockroachDB self-hosted use `cockroach start`
- These should be followed by the `locality` flag and any other flags

```
cockroach start --locality= \
region=us-east-1,zone=us-east-1b: \
region=us-east-1,zone=us-east-1a: \
region=us-east-1,zone=us-east-1c: \
region=us-east-2,zone=us-east-2b: \
region=us-east-2,zone=us-east-2a: \
region=us-east-2,zone=us-east-2c: \
region=us-west-1,zone=us-west-1b: \
region=us-west-1,zone=us-west-1a: \
region=us-west-1,zone=us-west-1c \
# ... other required flags go here
```

*Example using CockroachDB
Self-hosted*














CONFIGURING CLUSTER REGION USING COCKROACHDB DEDICATED

Regions & nodes

For highly-available single region clusters, choose a minimum of three nodes. Multi-region clusters must have a minimum of three regions and three nodes per region to survive zone and regional failures. [Read more](#) about our recommended multi-region configurations.

Region

Nodes

 N. Virginia (us-east4) 	3 	
 California (us-west2) 	3 	
 Montréal (northamerica-northeast1) 	3 	



**AFTER CONFIGURING CLUSTER REGIONS,
DATABASE REGIONS SHOULD BE CONFIGURED**



**DATABASES IN A MULTI-REGION CLUSTER THAT
HAVE NOT YET HAD THEIR PRIMARY REGIONS
SET WILL HAVE THEIR REPLICAS SPREAD AS
BROADLY AS POSSIBLE FOR RESILIENCY.**



MULTI-REGION DATABASES

- For database to be multi-region, it requires setting a primary region at minimum.
- When a primary region is added:
 - Tables will be rebalanced. All tables will have their voting replicas and leaseholders moved to the primary region by default.
 - Another name for this is *regional by table*.



MULTI-REGION DATABASES

- Once a database is configured as multi-region:
 - All data is stored within its assigned regions.
 - CockroachDB optimizes access to the database's data from the primary region.



CONFIGURING DATABASE REGIONS

- To add the first region, use the ALTER DATABASE ... PRIMARY REGION statement

```
ALTER DATABASE {database} SET PRIMARY REGION "us-east1";
```

- To add additional regions use the following syntax

```
ALTER database {database} ADD region "europe-west1";
```

- Regions can also be dropped from a database

```
ALTER DATABASE {database} DROP REGION "europe-west1";
```

- *The primary region can only be dropped if it's the last region remaining region*




**THE *TABLE LOCALITY* INDICATES HOW
COCKROACHDB OPTIMIZES ACCESS TO A
TABLE'S DATA IN A MULTI-REGION CLUSTER.**



GLOBAL TABLES



GLOBAL TABLE LOCALITY



- Tables optimized for fast reads from all regions.
- Global tables have slightly slower data write speeds.
- Global tables are useful in cases where data is read a lot more than it is written.



WHEN TO USE GLOBAL TABLES

- Reads cannot be stale for business or technical reasons (foreign key constraints).
- Your use case can accept higher write latency.
- Latency-sensitive reads cannot be tied to specific regions.



**GLOBAL TABLES ARE A GOOD CHOICE
FOR RARELY UPDATED LOOK-UP OR
REFERENCE TABLES.**



GLOBAL TABLE LOCALITY SYNTAX


```
ALTER TABLE {table} SET LOCALITY GLOBAL;
```



REGIONAL TABLES



REGIONAL TABLE LOCALITY



- Regional is the default locality for all multi-region database tables.
- A regional tables is optimized for fast reads and writes from its home region.
- Access to regional table data from other regions will be slower.



REGIONAL TABLE LOCALITY SYNTAX

- A regional table's home region defaults to the database's primary region. If you were to do this manually, the syntax would look like this:

```
ALTER TABLE {table} SET LOCALITY REGIONAL BY TABLE IN PRIMARY REGION;
```

- You can also set the locality without specifying the region. This will default again to the database's primary region.

```
ALTER TABLE {table} SET LOCALITY REGIONAL BY TABLE;
```

- To specify a non-default region, use the following syntax:

```
ALTER TABLE {table} SET LOCALITY REGIONAL BY TABLE IN "us-east-1";
```



REGIONAL BY TABLE LIMITATIONS

- Regional by table is the default and may not meet your needs if:
 - You need fast reads from all regions.
 - You need fast reads and writes from more than one region.

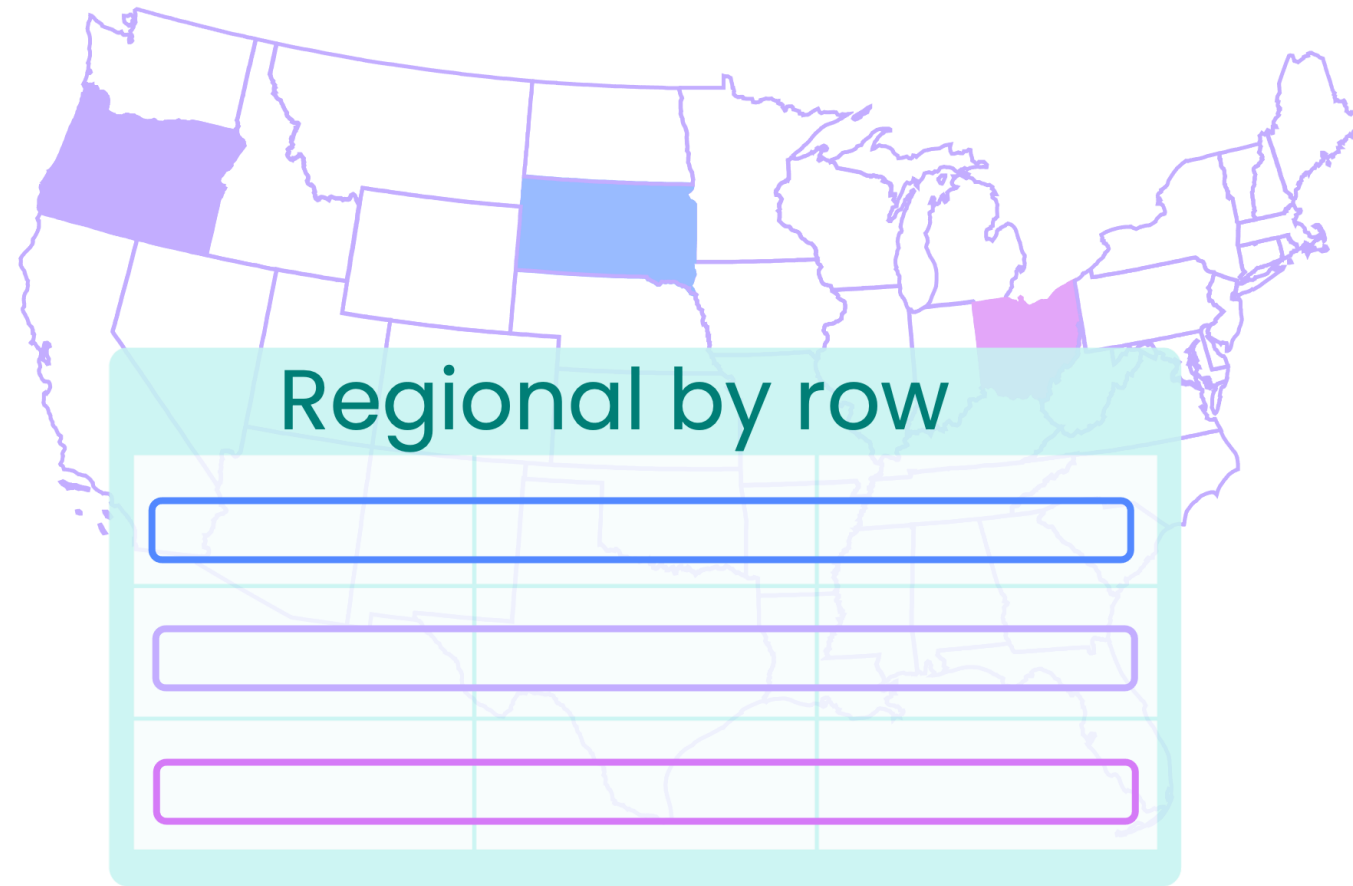


REGIONAL BY ROW TABLES


**WHEN YOU NEED FAST READS AND WRITES FROM MORE THAN
ONE REGION.**



REGIONAL BY ROW TABLE LOCALITY



- Assigns a region to each row
- Optimized for fast reads and writes from the region assigned to the row
- The table and its indexes are automatically partitioned on the region column.
- Each partition is optimized for access from a specific region.



REGIONAL BY ROW TABLE LOCALITY SYNTAX

- To make an existing table a regional by row table, use the following statement:

```
ALTER TABLE {table} SET LOCALITY REGIONAL BY ROW;
```

- This will create a hidden `crdb_region` column that represents the row's home region.
- You can use a column name other than `crdb_region` for the hidden column:

```
ALTER TABLE {table} SET LOCALITY REGIONAL BY ROW AS {column_name};
```



POPULATING THE `crdb_region` COLUMN

- For existing data, you can manually update the `crdb_region` column.

- Example:

```
UPDATE {table} SET crdb_region = 'us-west' WHERE city IN (...)
```


- For new inserts there are two options for populating the `crdb_region` field.
 - CockroachDB can set it automatically using the region of the gateway node (`gateway_region()`)
 - You can manually insert the data using the following:

```
INSERT INTO {table} (crdb_region, ...) VALUES ('us-east', ...);
```



SUMMARY - LATENCY AND LOCALITY

- CockroachDB multi-region allows you to address latency related to where your data is located.
- The concept of locality is present at the cluster, database and table level.
- A primary region must be added to a database to make it multi-region.



SUMMARY - TABLE LOCALITY

- *Global* table locality optimizes for fast read access from all regions with slightly slower write access.
- The default table locality, *regional*, optimizes for fast access from a single region.
- *Regional by row* table locality optimizes for fast access from a single region *specified at the row level*.

